

# Assgn3: Students' goals

Salvador Garcia, s1655274

10 February 2016

## 1 Description of the problem

229 students aged 7-13 from 9 schools were asked whether popularity or sporting ability was most important to them. The aim is to create an appropriate model that takes into account variability between the schools (if necessary).

**Statistical analysis.** Introduce the following random variables. Consider the  $i$ th student,  $i = 1, 2, \dots, N = 229$ , and set  $Z_i = 1$  if popularity is more important to the student than sporting ability, and  $Z_i = 0$  otherwise.

## 2 case 1: $\theta_1 = \theta_2 = \dots = \theta_9$

## 3 case 2: with independent non-informative priors

## 4 case 3: with a hierarchical exchangeable priors

The data have the following distribution:

$Z_i | \theta_k \sim \text{Bern}(\theta_k)$ , where  $k$  is the number of the school the  $i$ th student is from, independently given  $\theta_k$ ,  $\theta_k \in (0, 1)$ ,  $k = 1, 2, \dots, 9$ .

The purpose of this assignment is to determine if there is a difference between  $\mu_1$  and  $\mu_2$ . In terms of the problem, to find if there is a relevant difference of gene expression in the two different tissues. In traditional statistics, this can be addressed by testing the hypothesis  $\mu_1 = \mu_2$  or, equivalently,  $\mu_1 - \mu_2 = 0$  given some  $\alpha$ . In Bayesian statistics something similar can be done, but it is in terms of the posterior distribution of  $\delta = \mu_1 - \mu_2$ .

## 5 Prior distribution

### 5.1 Priors for $\mu_i$

In this example we are working with normal distributions (that have two parameters) of two populations,  $\mu_1$  and  $\sigma_1$  for the first one and  $\mu_2$  and  $\sigma_2$  for the second one. In order to address the information of the  $\mu_i$  parameters, two uninformative priors for  $\mu_1$  and  $\mu_2$  will be used:  $\mu_i \sim N(0, A^2)$  with  $A$  very big. This way, the prior distribution of  $\mu_1$  and  $\mu_2$  will be as the fig. 1. It is easy to see that this distribution assigns almost evenly the probabilities to all values from  $(-\infty, \infty)$ . Then, can be viewed as non-informative (The idea of apriori with a very big variance is important in this assignment).

Now, given the priors of  $\mu_1$  and  $\mu_2$ , is necessary to think about the priors of  $\sigma_1$  and  $\sigma_2$ . Is important to first give the priors  $\mu_i$  because  $\sigma_i$  depends on its value.

### 5.2 Priors for $\sigma_i$

For  $\sigma_1$  and  $\sigma_2$  a gamma distribution with parameters  $\text{shape}(\alpha) = .001$  and  $\text{rate}(\beta) = .001$  will be used. This distribution is shown in fig. 1. Remembering, the mean of the gamma distribution is  $\frac{\alpha}{\beta}$ , and the variance is  $\frac{\alpha}{\beta^2}$ . Then this distribution have  $\text{mean} = 1$  and  $\text{variance} = 1000$ .

Now, information from a previous study is used as a new apriori for  $\sigma_i$ . For the first one, the study says that the average precision (inverse of variance) is 11.11 for the cases and 100 for the control, but the variability is not known. This information given is presented below eq. (1).




Figure 1: Uninformative normal and gamma distributions

$$\frac{\alpha_1}{\beta_1} = 11.11 \quad , \quad \frac{\alpha_2}{\beta_2} = 100 \quad (1)$$

Now, the variability of these values it is not known, so two assumptions to get two different scenarios will be made. The first one is that the prior variance is 10 and the other that the prior variance is 1000. Then, for both scenarios, this information can be interpreted as followseq. (2):

$$\frac{\alpha_i}{\beta_i^2} = 10 \quad , \quad \frac{\alpha_i}{\beta_i^2} = 1000 \quad (2)$$

With these formulas, a system of equations can be created to find  $\alpha_1$ ,  $\alpha_2$ ,  $\beta_1$ , and  $\beta_2$ . Using simple algebra, is easy to find that the parameters for the large variance scenario are eq. (3) and for the low variance scenario are eq. (4).

$$\alpha_1 = \frac{11.11^2}{1000} \quad , \quad \beta_1 = \frac{11.11}{1000} \quad , \quad \alpha_2 = \frac{100^2}{1000} \quad , \quad \beta_2 = \frac{100}{1000} \quad (3)$$

$$\alpha_1 = \frac{11.11^2}{10} \quad , \quad \beta_1 = \frac{11.11}{10} \quad , \quad \alpha_2 = \frac{100^2}{10} \quad , \quad \beta_2 = \frac{100}{10} \quad (4)$$

This two *a priori* distributions are used along with the non-informative distribution selected. As the hint says, it is important to distinguish the difference between precision and variance in the model. The one that is modelled though a gamma distribution is the precision. So the variability (variance) of 10 and 1000 shown in this section is the variability of the gamma distribution (the prior) of the precision.

## 6 Posterior inference

In this section the posterior analysis of  $(\sigma_1, \sigma_2)$  will be given. This is done with the relations  $\sigma_1 = \frac{1}{\tau_1}$  and  $\sigma_2 = \frac{1}{\tau_2}$ . For this analysis, two chains will be used with different initial points (due to the convergence property, at the end both chains will end in the same value).

### 6.1 Posterior distributions

#### 6.1.1 non informative prior

Table 1: Summary of analysis with non informative prior

parameters	mean	sd	2.5%	50%	97.5%
$\mu_1$	4.03	0.0119	4.007	4.03	4.053
$\mu_2$	2.59	0.005	2.58	2.59	2.6
$\delta$	1.44	0.0129	1.415	1.44	1.465
$\sigma_1$	0.2907	0.0084	0.2749	0.2905	0.3081
$\sigma_2$	0.1105	0.0035	0.1038	0.1104	0.1175

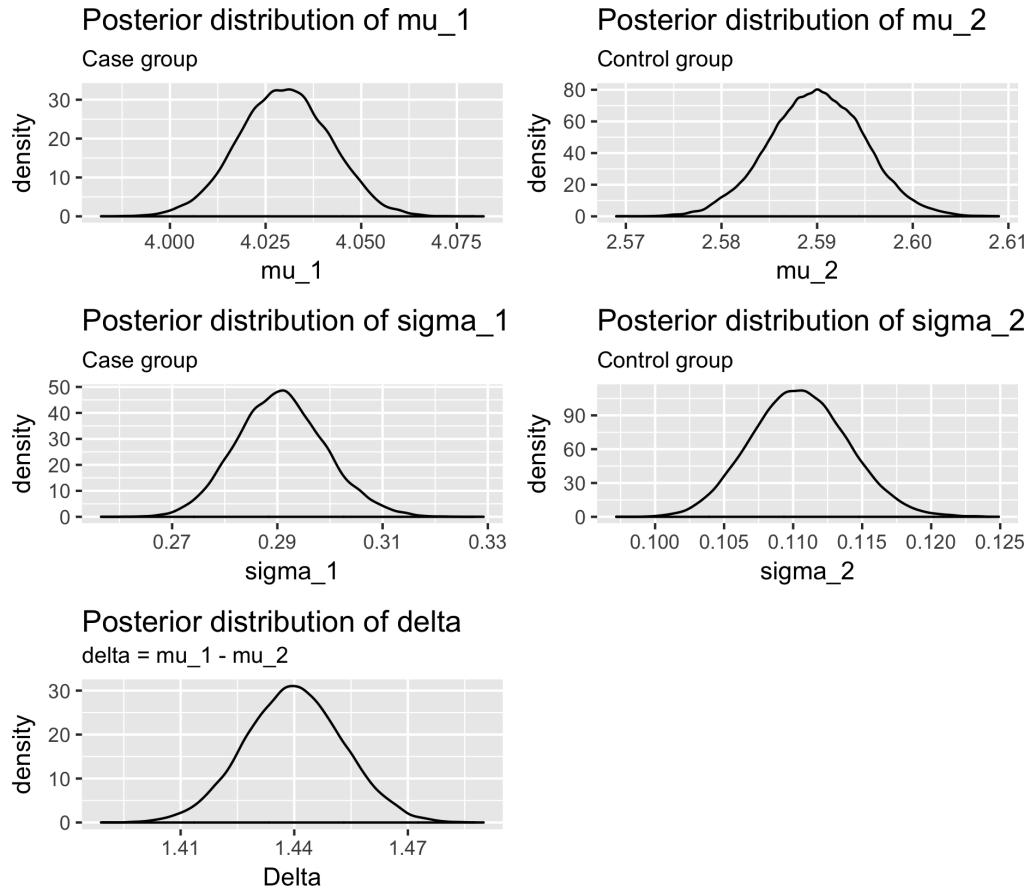


Figure 2: Posterior distributions with uninformative prior

In this example, as the fig. 2 and table 1 show, there is an important difference between  $\mu_1$  and  $\mu_2$ . In fact, the credibility intervals (95%) of each are (4.007, 4.053) and (2.58, 2.6) and equivalently for  $\delta$ : (1.415, 1.465). With this prior, we can conclude that there is a difference between  $\mu_1$  and  $\mu_2$ .

### 6.1.2 Informative prior with low variance

Table 2: Summary of analysis with informative prior (low variance)

parameters	mean	sd	2.5%	50%	97.5%
$\mu_1$	4.0299	0.0119	4.007	4.03	4.053
$\mu_2$	2.59	0.0046	2.581	2.59	2.599
$\delta$	1.4399	0.0128	1.415	1.44	1.465
$\sigma_1$	0.291	0.0082	0.2755	0.2909	0.3081
$\sigma_2$	0.1021	0.0014	0.0994	0.1021	0.105

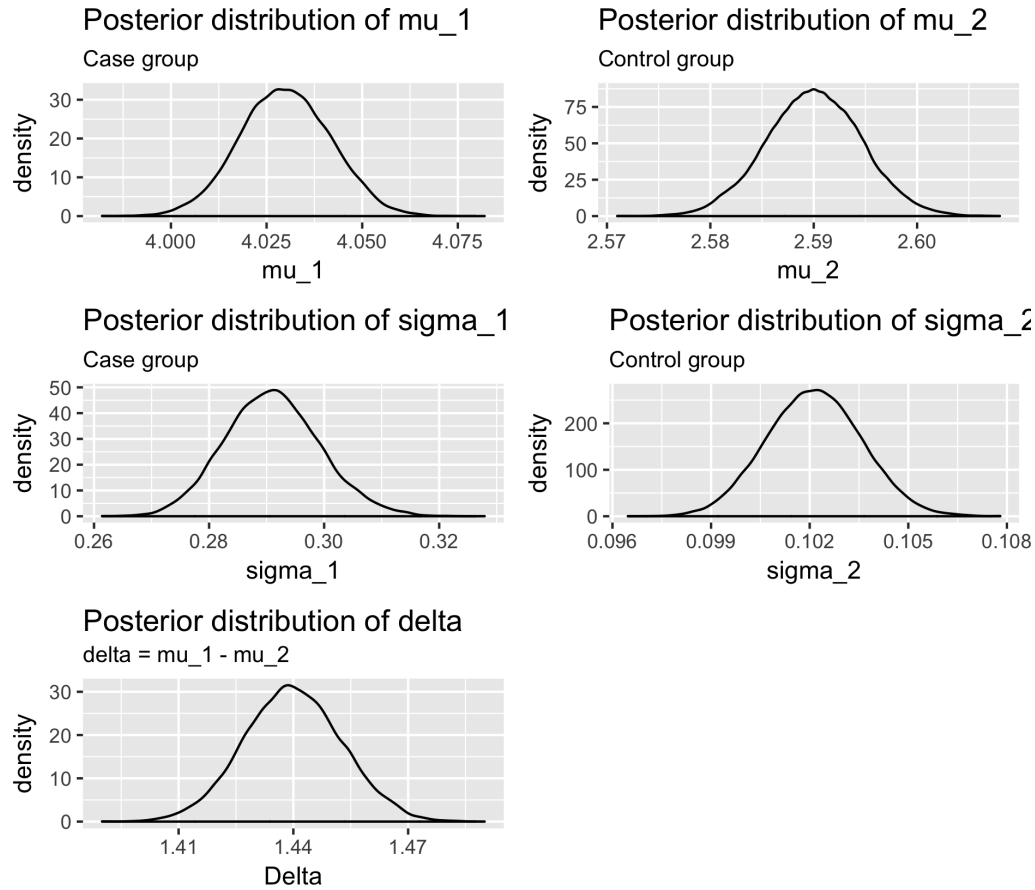


Figure 3: Posterior distributions with informative prior (low variance)

In this example, as the fig. 3 and table 2 show, there is an important difference between  $\mu_1$  and  $\mu_2$ . In fact, the credibility intervals (95%) of each are (4.007, 4.053) and (2.581, 2.599) and equivalently for  $\delta$ : (1.415, 1.465). The conclusions are the same than the previous example (with the non informative prior). But is important to distinguish that the posterior distribution of  $\sigma_1$  and  $\sigma_2$  are a little bit different. (0.2907 vs 0.291 and 0.1105 vs 0.1021).

### 6.1.3 Informative prior with large variance

Table 3: Summary of analysis with informative prior (large variance)

parameters	mean	sd	2.5%	50%	97.5%
$\mu_1$	4.0299	0.0119	4.007	4.03	4.053
$\mu_2$	2.59	0.005	2.58	2.59	2.6
$\delta$	1.4399	0.0129	1.415	1.44	1.465
$\sigma_1$	0.2907	0.0084	0.2748	0.2905	0.3081
$\sigma_2$	0.1099	0.0034	0.1034	0.1099	0.1168

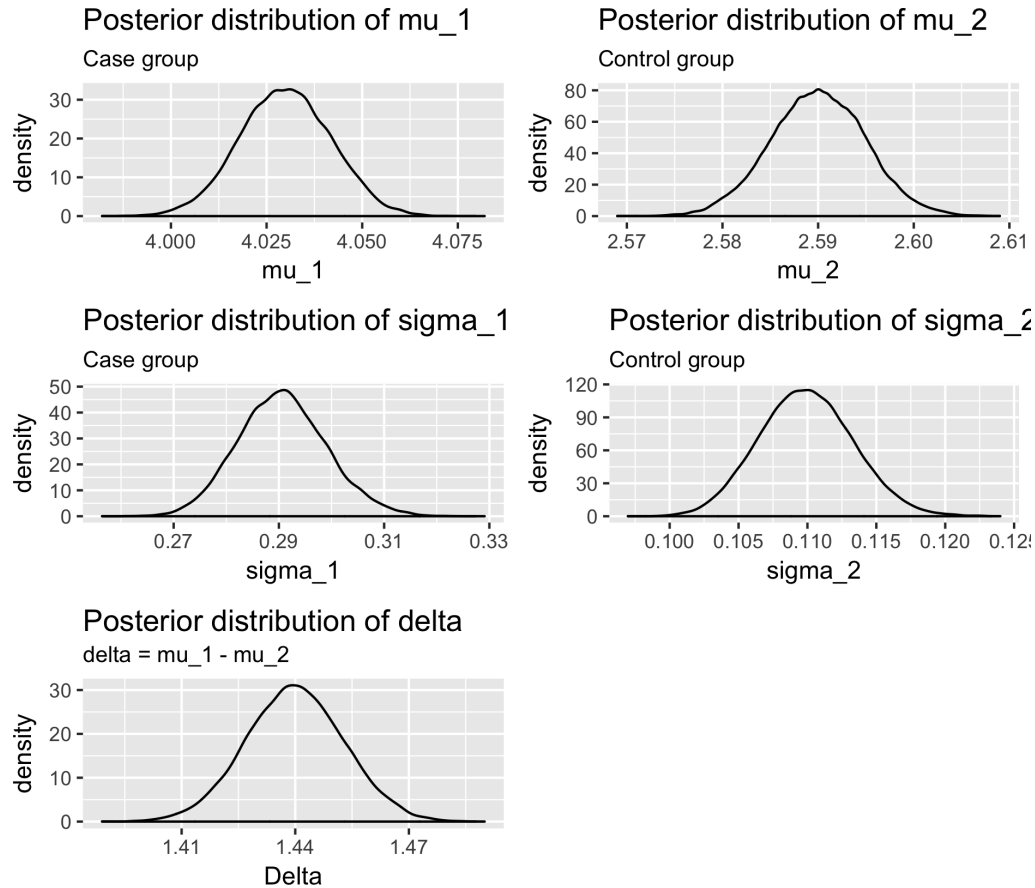


Figure 4: Posterior distributions with informative prior (large variance)

At last, the a informative a priori with large variance is used. The conclusions are the same than the previous two examples, but the estimations of the parameters and the standards deviation of them are almost the same with the first example (non informative). The differences are of the magnitude of 0.0001.

## 6.2 Posterior convergence

For the posterior distributions showed in the last subsections, two chains were used with the initial values:

$$\begin{aligned} \mu_1 = 10 \quad , \quad \mu_2 = 0 \quad , \quad \tau_1 = 0.1 \quad , \quad \tau_2 = 0.5 \\ \mu_1 = -10 \quad , \quad \mu_2 = -10 \quad , \quad \tau_1 = 0.2 \quad , \quad \tau_2 = 0.2 \end{aligned} \quad (5)$$

### 6.2.1 non informative prior

In order to determine if the chain have converged, The Gelman-Rubin statistics was tried. For the four parameters there is a good convergence, but for  $\sigma_2$  maybe more iterations would be better. (There are little differences between the chains, but it is a very good convergence). Another analysis made was to check the autocorrelations coefficients. The first blue line is 1 because the lag is 0, so is the correlation between the same element. In other parts of the chain seems that there are no important correlation between the consecutive points (at least the shown).

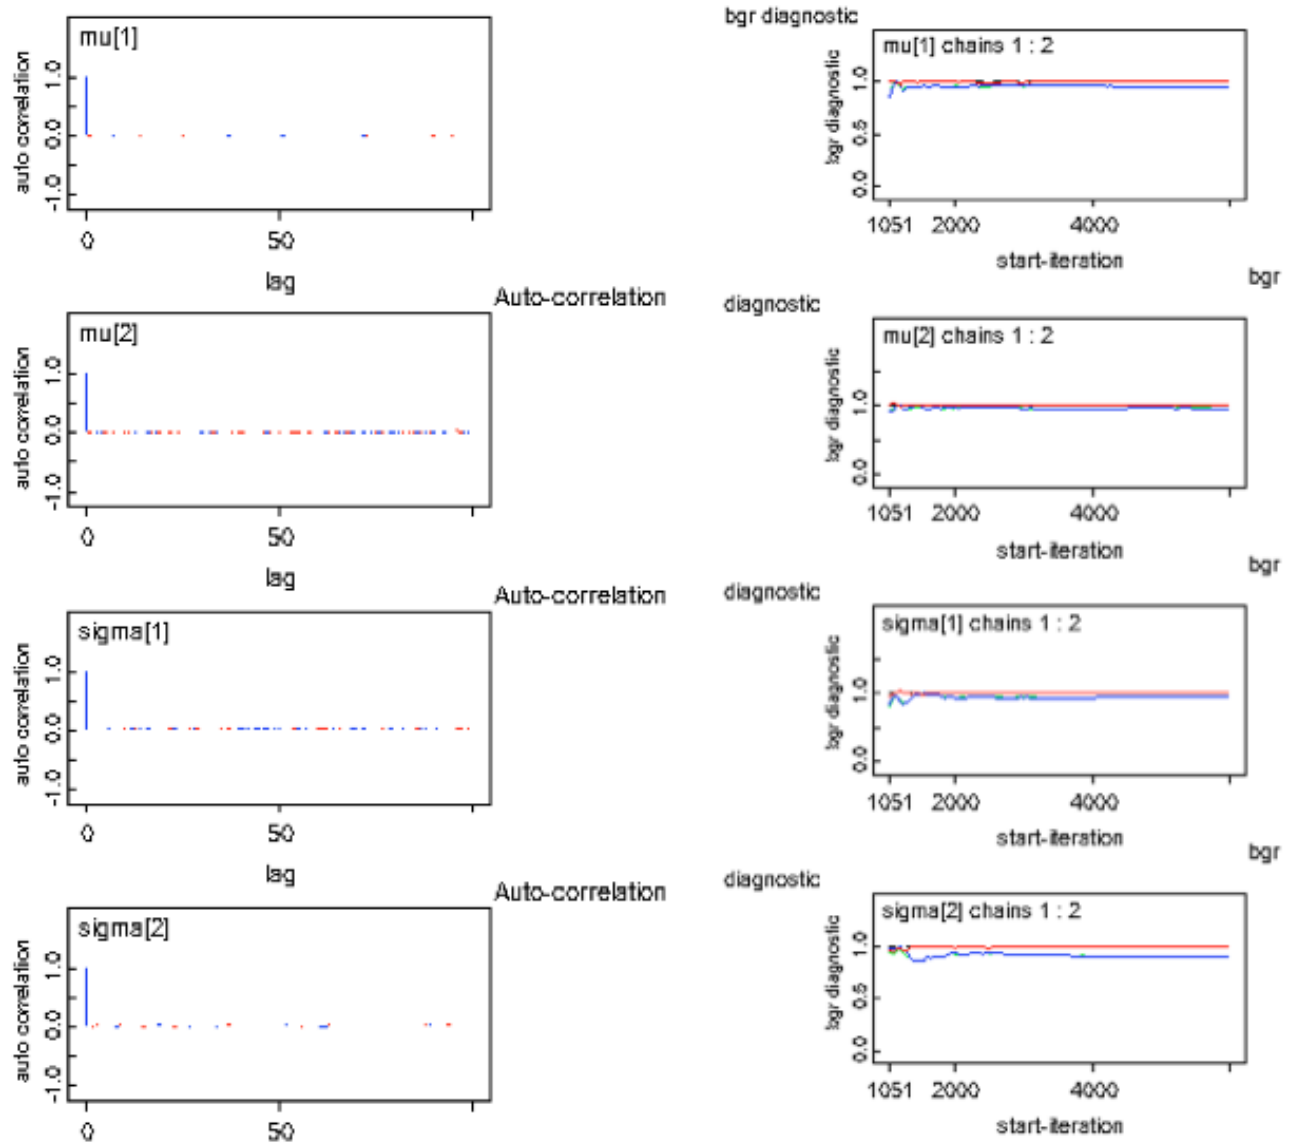


Figure 5: Autocorrelation and Gelman-Rubin statistic (non informative prior)

Now, the history of the parameters is shown in fig. 6. Is easy to see that the chain is exploring the same area and there are almost any point outside the central area of the distributions. So, it looks very good the convergence of the chain. The burning iterations were 1500. Before that number, the chain moves trying to explore the parameters domain (different initial points should converge to the same point). This is an indication that the optimum is still not achieved, and if these iterations are considered, the analysis will be biased. The autocorrelation analysis is very important. As the MCMC name states, is a Markov Chain, so the position on the next iterations depends directly on the position in the current iteration. It is important to thin in order to de-correlate each iteration (although some authors says that most of the times it is not necessary to thinning). The MC error is good when checking the fig. 7. According with the BIAS project, the MC error compares with the SD of the posterior distribution in a fraction  $\frac{SD}{\sqrt{(iterations)}}$  also that the thumb error is that should be less than 1% of the posterior SD is a good value. In all the estimates this is achieved.




Figure 6: History of parameters (non informative)

				Node statistics			start	sample
		mean	sd	MC_error	val2.5pc	median	val97.5pc	
delta		1.44	0.0129	9.063E-5	1.415	1.44	1.465	1001
mu[1]		4.03	0.01186	8.497E-5	4.007	4.03	4.053	1001
mu[2]		2.59	0.004976	3.549E-5	2.58	2.59	2.6	1001
sigma[1]		0.2907	0.008403	6.414E-5	0.2748	0.2905	0.3081	1001
sigma[2]		0.1105	0.0035	2.34E-5	0.1038	0.1104	0.1175	1001
tau[1]		11.86	0.6841	0.005221	10.53	11.85	13.24	1001
tau[2]		82.19	5.201	0.03486	72.41	82.03	92.75	1001

Figure 7: Summary of model (non informative)

### 6.3 Informative prior with low variance

The Gelman-Rubin statistics was tried for this example too. For the four parameters there is a good convergence, but again,  $\sigma_2$  could have better performance if more iterations are added. (There are little differences between the chains, but it is a very good convergence). The autocorrelation coefficients also were calculated. The conclusions are the same than the previous example.

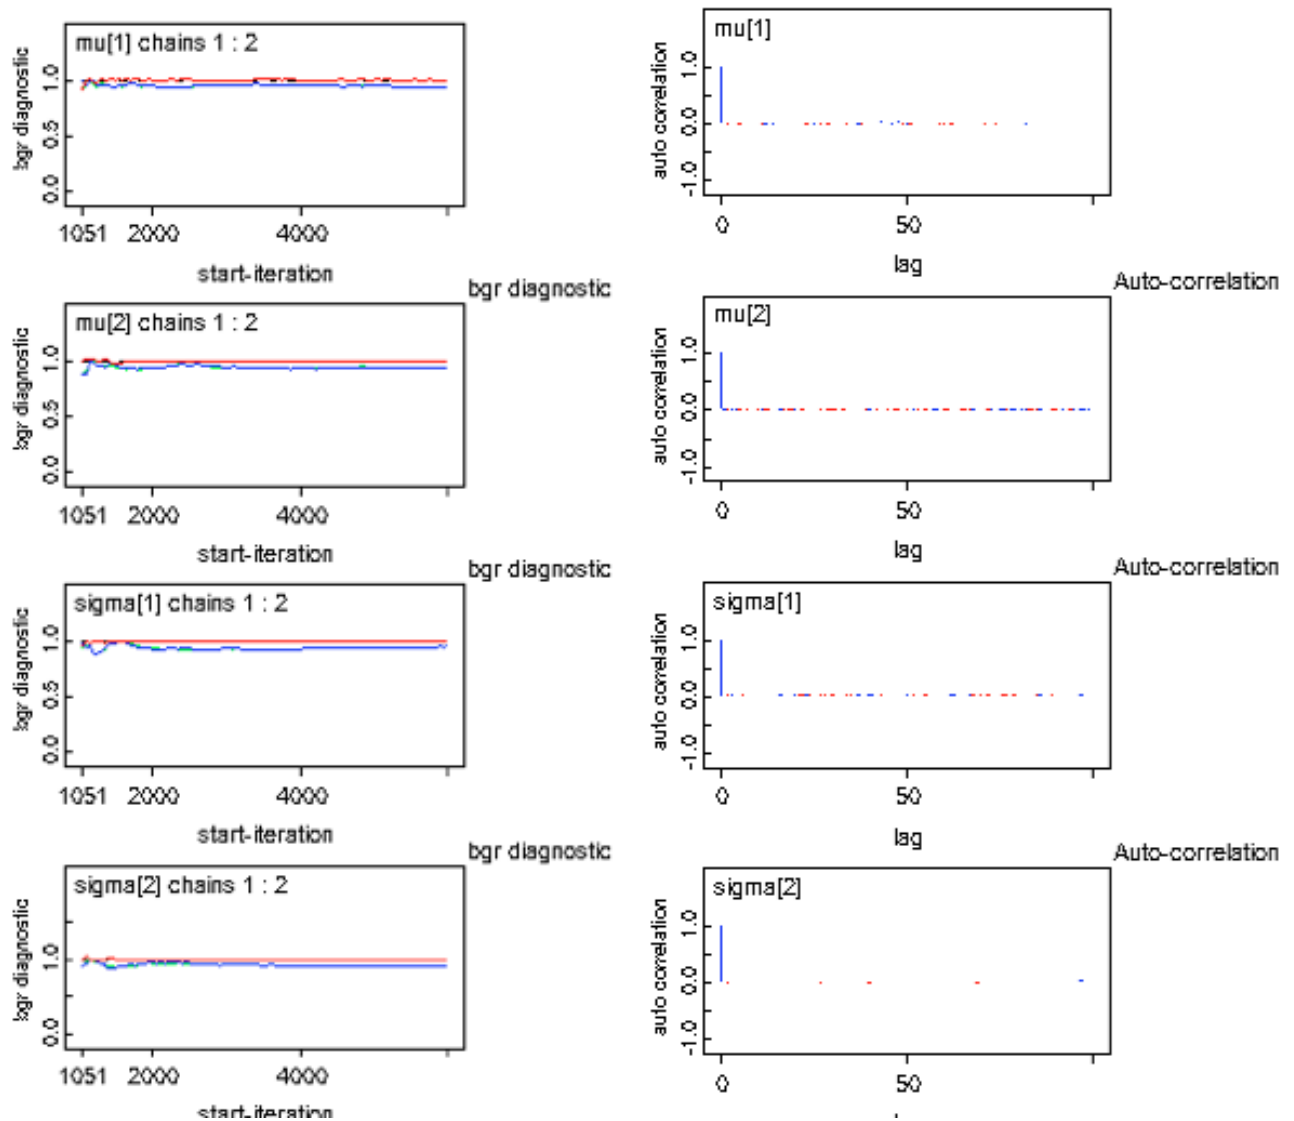


Figure 8: Autocorrelation and Gelman-Rubin statistic (low variance prior)



Again, the history of the parameters is shown in fig. 9. The chains seems to have converged and there is no problem associated with convergence. In this example the burning iterations were 1500 too. The values of the MC error are also below the 1% thumb rule.

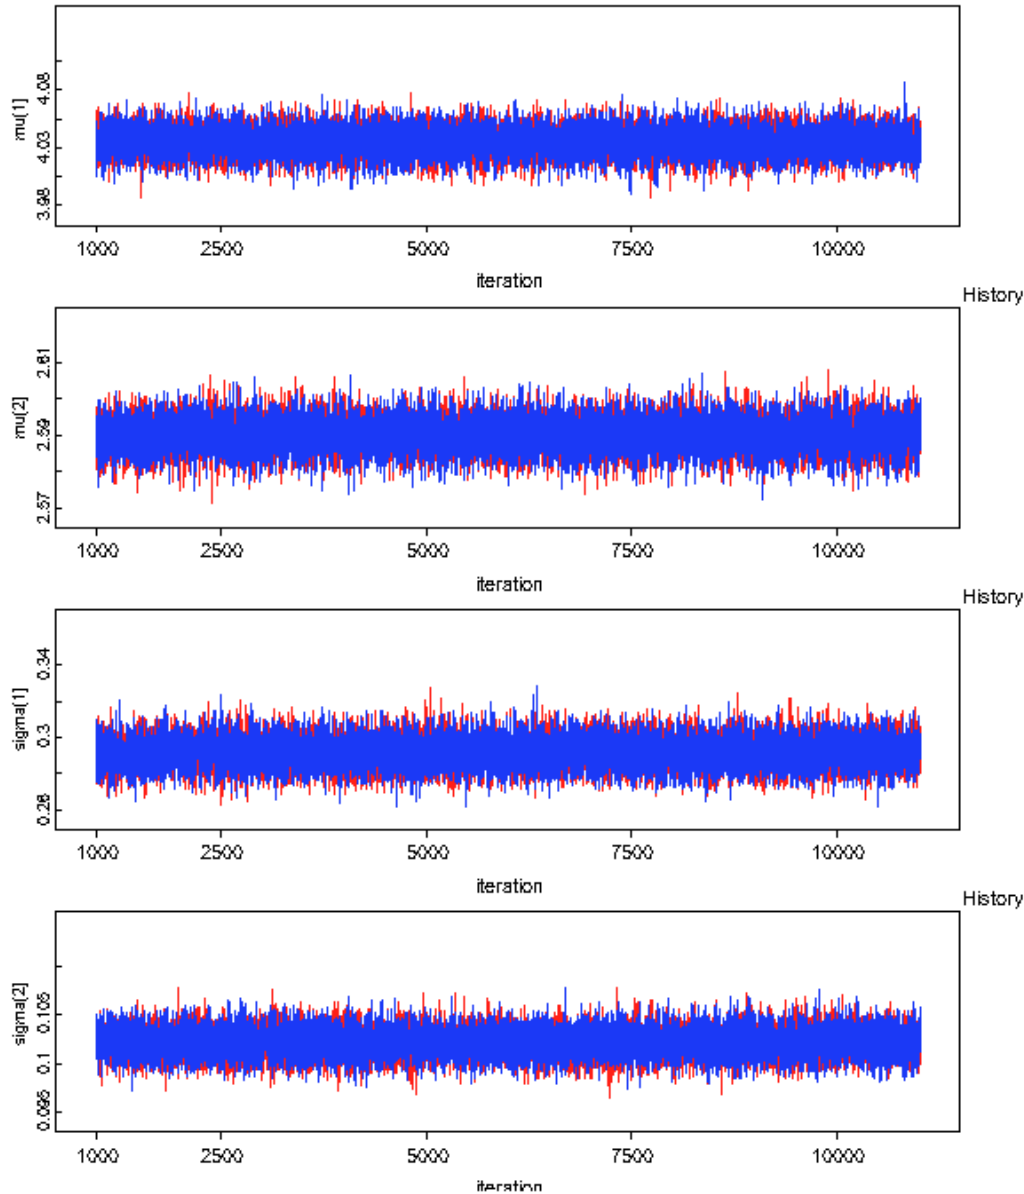


Figure 9: History of parameters (low variance prior)

	mean	sd	MC_error	val2.5pc	median	val97.5pc	start	sample
delta	1.44	0.01274	8.634E-5	1.415	1.44	1.465	1001	20000
mu[1]	4.03	0.01184	7.799E-5	4.007	4.03	4.053	1001	20000
mu[2]	2.59	0.004595	3.354E-5	2.581	2.59	2.599	1001	20000
sigma[1]	0.291	0.008238	5.756E-5	0.2755	0.2909	0.308	1001	20000
sigma[2]	0.1021	0.001446	9.963E-6	0.09937	0.1021	0.105	1001	20000
tau[1]	11.84	0.6684	0.004674	10.54	11.82	13.18	1001	20000
tau[2]	95.91	2.715	0.0187	90.73	95.86	101.3	1001	20000

Figure 10: Summary of model (low variance prior)

## 6.4 Informative prior with large variance

The Gelman-Rubin statistics was tried for this example too. For the four parameters there is a good convergence, but again,  $\sigma_2$  could have better performance if more iterations are added. (There are little differences between the chains, but it is a very good convergence). The autocorrelation coefficients also were calculated. The conclusions are the same than the previous two examples.

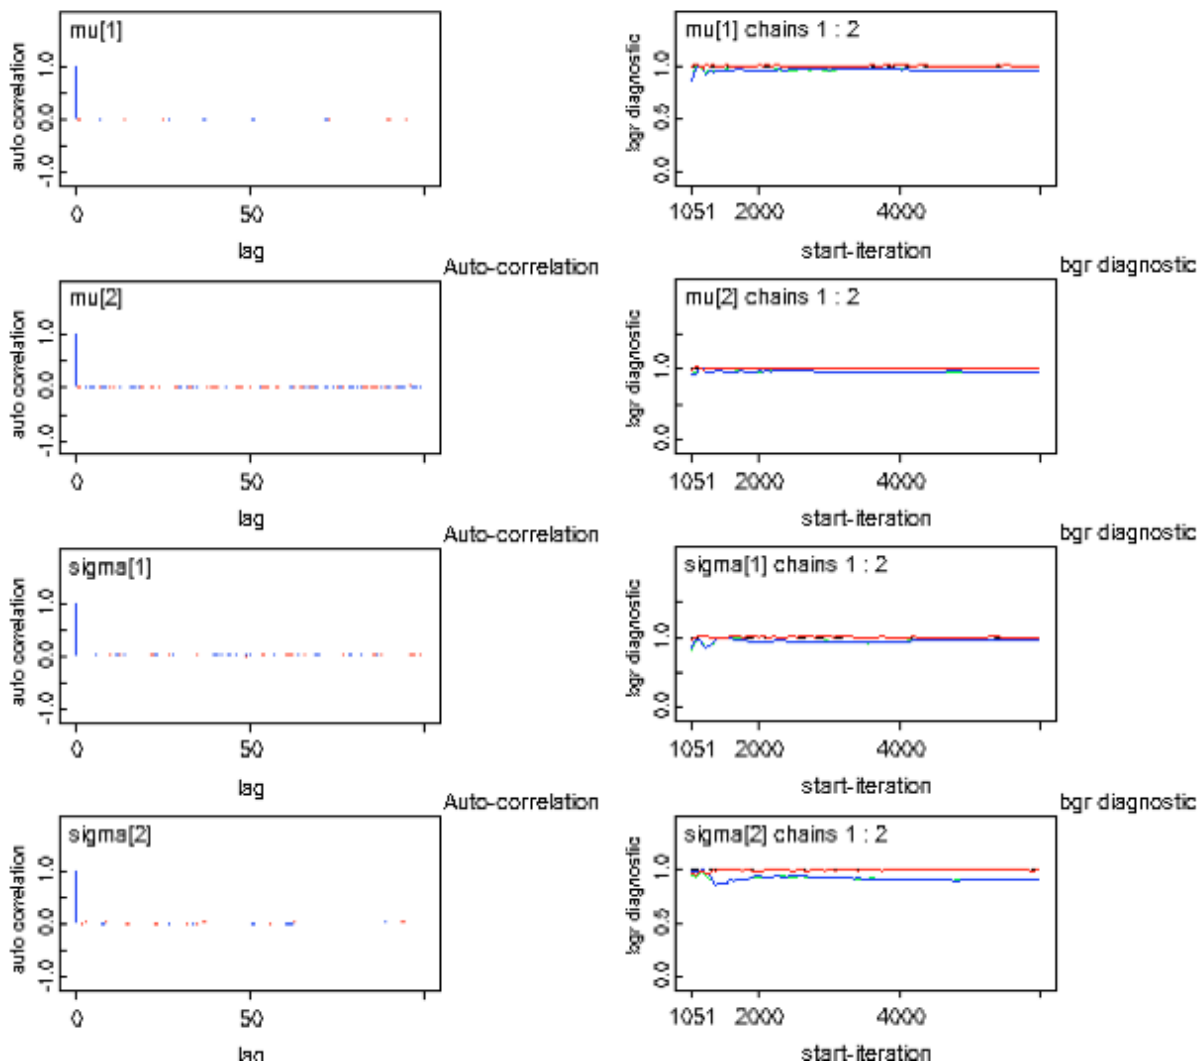


Figure 11: Autocorrelation and Gelman-Rubin statistic (large variance prior)

Again, the history of the parameters is shown in fig. 12. The chains seems to have converged and there is no problem associated with convergence. In this example the burning iterations were 1500 too. The values of the MC error are also below the 1% thumb rule.




Figure 12: History of parameters (large variance prior)

	mean	sd	MC_error	val2.5pc	median	val97.5pc	start	sample
delta	1.44	0.0129	9.063E-5	1.415	1.44	1.465	1001	20000
mu[1]	4.03	0.01186	8.497E-5	4.007	4.03	4.053	1001	20000
mu[2]	2.59	0.004976	3.549E-5	2.58	2.59	2.6	1001	20000
sigma[1]	0.2907	0.008403	6.414E-5	0.2748	0.2905	0.3081	1001	20000
sigma[2]	0.1105	0.0035	2.34E-5	0.1038	0.1104	0.1175	1001	20000
tau[1]	11.86	0.6841	0.005221	10.53	11.85	13.24	1001	20000
tau[2]	82.19	5.201	0.03486	72.41	82.03	92.75	1001	20000

Figure 13: Summary of model (large variance prior)

## 7 Model comparison

In this assignment three prior distributions were used: the first was an uninformative prior, while the second and third were informative ones. For these two, one was a distribution with low variance and the other one was a distribution with large variance. The intuition behind the apriori non informative is that this has a very large variance, so the model does not restrict the parameter into some space in the domain of the parameter. This way, a normal distribution with a very big variance can be used as a non informative prior. Now, the informative distribution with large variance is expected to behave similar than the non informative, as can be see in the fig. 14. Also, in this graph is easy to how is similar to the informative distribution with small variance (except in the  $\sigma_2$  parameter).

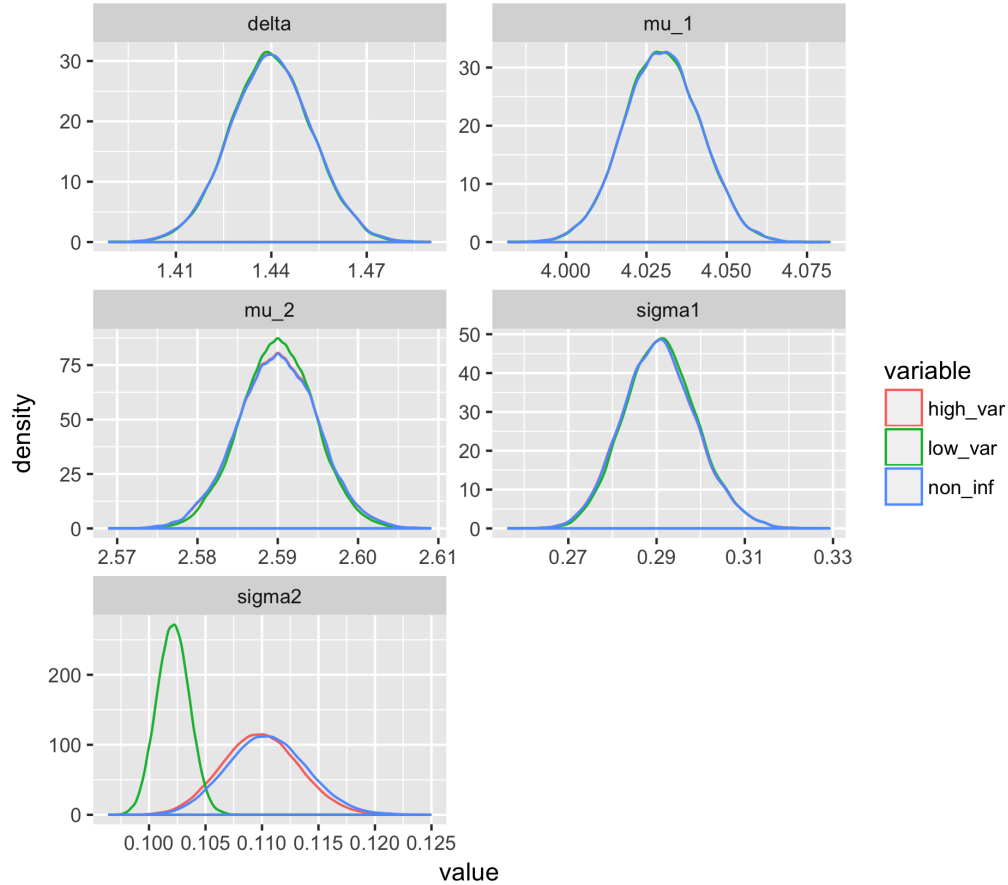


Figure 14: Comparison of posterior distributions with different apriori

Although there is a difference in  $\sigma_2$ , when the scale is observed, then this can be negligible (compared with the original a priori which has a mean of 100).

In summary, the specification of a prior far from the likelihood and also with small variance, can affect the posterior distribution. But in this example it just change the mean of the distributions just by a small amount. It is important to say that the modifications of the variance can have more important changes in the posterior distribution.

## A R graphs code

For the plots in R, the package tidyverse and R2OpenBUGS were used. Also, the package wine was used. BUGS models were sent via email

```
#Load the OpenBUGS Package – make sure XQuartz is running
library(tidyverse)
library(R2OpenBUGS)

#define the model
nummodel <- function(){
  # Model
  # distribution of the sample means
  barx ~ dnorm(mu[1], tau.mean[1])
  bary ~ dnorm(mu[2], tau.mean[2])
  # distribution of the sample variances
  s2x ~ dgamma(n1, tau1.2)
  s2y ~ dgamma(n2, tau2.2)
  # Priors
  for (j in 1 : 2){
    mu[j] ~ dnorm(0, 0.000001)
    # standard deviation of error distribution
    sigma[j] <- sqrt(1 / tau[j])
  }
  tau[1] ~ dgamma(0.01,0.01)
  tau[2] ~ dgamma(0.01,0.01)
  # introduce the difference between the means (cases–controls)
  delta <- mu[1] – mu[2]
  # introduce the precisions of the means
  tau.mean[1] <- tau[1]*n*N
  tau.mean[2] <- tau[2]*m*N
  # introduce the scale of the gamma distribution (chi–square)
  tau1.2 <- n*N* tau[1]/2
  tau2.2 <- m*N*tau[2]/2
  # introduce the shape of the gamma distribution (chi–square)
  n1 <- (n*N–1)/2
  n2 <- (m*N–1)/2
}

# write the model code out to a file
write.model(nummodel, "nummodel.txt")
model.file1 = paste(getwd(),"nummodel.txt", sep="/")
## and let's take a look:
file.show("nummodel.txt")

#prepare the data for input into OpenBUGS
n = 12
m=10
N = 50
barx = 4.03
bary = 2.59
s2x = 0.0841
s2y = 0.0121
data <- list ("n", "m", "N", "barx", "bary", "s2x", "s2y")
```

```

# initialization of variables
inits <- function(){
  list (chain1 <- c("mu[1]=10", "mu[2] = 0", "tau[1] = 0.1", "tau[2] = 0.5"),
        chain2 <- c("mu[1]=-10", "mu[2] = -10", "tau[1] = 0.2", "tau[2] = 0.2"))}

#set the WINE working directory and the directory to OpenBUGS – change the OpenBUGS.exe location as
  necessary
WINE="/usr/local/bin/wine"
WINEPATH="/usr/local/bin/winepath"
OpenBUGS.pgm="/Users/salvadorgarcia/.wine/drive_c/Program
  Files/OpenBUGS/OpenBUGS323/OpenBUGS.exe"

#these are the parameters to save
parameters = c("mu", "tau", "delta", "sigma")

#run the model
schools.sim <- bugs(data, inits,
                    model.file = model.file1,
                    parameters=parameters,
                    n.chains = 2,
                    n.iter = 11000,
                    OpenBUGS.pgm=OpenBUGS.pgm,
                    WINE=WINE,
                    WINEPATH=WINEPATH,
                    useWINE=T,
                    n.burnin = 1500)

schools.sim$summary %>% round(4) %>% write.table(pipe("pbcopy"))

modelo_noninf <- schools.sim
p1 <- schools.sim$sims.list$delta %>%
  data.frame() %>% setNames(c("Delta")) %>%
  ggplot(aes(x= Delta))+
  geom_density() +
  labs( title = "Posterior distribution of delta", subtitle = "delta = mu_1 - mu_2")

p2 <- schools.sim$sims.list$mu[,1] %>%
  data.frame() %>% setNames(c("mu_1")) %>%
  ggplot(aes(x= mu_1))+
  geom_density() +
  labs( title = "Posterior distribution of mu_1", subtitle = "Case group")

p3 <- schools.sim$sims.list$mu[,2] %>%
  data.frame() %>% setNames(c("mu_2")) %>%
  ggplot(aes(x= mu_2))+
  geom_density() +
  labs( title = "Posterior distribution of mu_2", subtitle = "Control group")

p4 <- schools.sim$sims.list$sigma[,1] %>%
  data.frame() %>% setNames(c("sigma_1")) %>%
  ggplot(aes(x= sigma_1))+

```

```

geom_density() +
labs( title = "Posterior distribution of sigma_1", subtitle = "Case group")

p5 <- schools.sim$sims.list$sigma[,2] %>%
  data.frame() %>% setNames(c("sigma_2")) %>%
  ggplot(aes(x= sigma_2))+
  geom_density() +
  labs( title = "Posterior distribution of sigma_2", subtitle = "Control group")

ggsave(plot = grid.arrange(p2,p3,p4,p5,p1, ncol = 2),
  filename = "../12_Assignment2/imgs/01_posterior_pars.png",
  width = 16, height = 14, units = "cm")

```

---

Code for the comparison

---

```

# comparison

# mu[1]
mu_1 <- data.frame(high_var = modelo_high$sims.list$mu[,1],
  low_var = modelo_low$sims.list$mu[,1],
  non_inf = modelo_noninf$sims.list$mu[,1]) %>%
  mutate(par = "mu_1")
# mu[2]
mu_2 <- data.frame(high_var = modelo_high$sims.list$mu[,2],
  low_var = modelo_low$sims.list$mu[,2],
  non_inf = modelo_noninf$sims.list$mu[,2]) %>%
  mutate(par = "mu_2")
# sigma[1]
sigma_1 <- data.frame(high_var = modelo_high$sims.list$sigma[,1],
  low_var = modelo_low$sims.list$sigma[,1],
  non_inf = modelo_noninf$sims.list$sigma[,1]) %>%
  mutate(par = "sigma1")

# sigma[2]
sigma_2 <- data.frame(high_var = modelo_high$sims.list$sigma[,2],
  low_var = modelo_low$sims.list$sigma[,2],
  non_inf = modelo_noninf$sims.list$sigma[,2]) %>%
  mutate(par = "sigma2")
# delta
delta <- data.frame(high_var = modelo_high$sims.list$delta,
  low_var = modelo_low$sims.list$delta,
  non_inf = modelo_noninf$sims.list$delta) %>%
  mutate(par = "delta")

pars_df <- rbind(mu_1, mu_2, sigma_1, sigma_2, delta)
p1 <- pars_df %>% gather(variable, value, high_var:non_inf) %>%
  ggplot(aes(x = value, color = variable)) +
  geom_density() +
  facet_wrap(~par, scale = "free", ncol = 2)

ggsave(plot = p1,
  filename = "../12_Assignment2/imgs/comparison.png",
  width = 16, height = 14, units = "cm")

```